Om Prakash Gunja

AI/ML Engineer

IL | (872)- 218 - 5054 | gunjaomprakash@gmail.com

SUMMARY

AI/ML Engineer with 3+ years of experience designing and deploying advanced machine learning and deep learning solutions. Proficient in Python, SQL, TensorFlow, PyTorch, and Scikit-learn, with expertise in CNNs, RNNs, LSTMs, GANs, and Transformers for NLP and computer vision applications. Skilled in big data technologies (Hadoop, Spark, Kafka) and cloud platforms (AWS, Azure, Google Cloud), leveraging Agile and Waterfall methodologies to deliver scalable AI solutions. Adept at using Pandas, NumPy, Matplotlib, and Seaborn for data analysis and visualization and implementing DevOps practices with Docker, Kubernetes, Git, and Jenkins to streamline model deployment. Recognized for driving operational efficiency and delivering impactful AI-driven insights in enterprise and IoT environments.

SKILLS

Methodology: SDLC, Agile, Waterfall
Programming Languages: Python, R, Java, SQL
Machine Learning Frameworks: TensorFlow, PyTorch, Scikit-learn, Keras
Deep Learning: CNNs, RNNs, LSTMs, GANs, Transformers
Data Science Tools: Pandas, NumPy, Matplotlib, Seaborn
Big Data Technologies: Hadoop, Spark, Apache Kafka
Cloud Platforms: AWS (S3, EC2, SageMaker), Azure, Google Cloud
DevOps Tools: Docker, Kubernetes, Git, Jenkins
Natural Language Processing & Generative AI: NLP, Hugging Face Transformers, LangChain, Prompt Engineering, Retrieval-Augmented Generation, Vector Databases, LLM Fine-Tuning, Chatbot Development, Generative AI Applications
Computer Vision & Visualization: OpenCV, Data Visualization (Tableau, Power BI, Seaborn, Matplotlib)

EXPERIENCE

ServiceNow, IL | Nov 2024 – Current | AI/ML Engineer

- Developed end-to-end machine learning pipelines using Python, Pandas, and Scikit-learn to preprocess and analyze large-scale enterprise datasets, improving predictive model accuracy by 15% for IT service management workflows.
- Implemented Convolutional Neural Networks (CNNs) with TensorFlow to enhance image-based anomaly detection in IT infrastructure monitoring, reducing false positives by 20%.
- Led Agile-based project sprints to design and deploy Natural Language Processing (NLP) models for automated ticket classification, streamlining issue resolution processes by 25%.
- Utilized AWS SageMaker to train and deploy scalable machine learning models, optimizing resource allocation for cloud-based IT service analytics.
- Created data visualizations using Matplotlib and Seaborn to present model performance metrics to stakeholders, enabling data-driven decision-making for service optimization.
- Integrated Git for version control in collaborative ML model development, ensuring seamless code reviews and maintaining project integrity across distributed teams.
- Designed and executed SQL-based data extraction pipelines to aggregate structured data from ServiceNow's IT service databases, improving data accessibility for model training.

Neon IT Systems, India | Jan 2020 - Dec 2022 | AI/ML Engineer

- Built and optimized Recurrent Neural Networks (RNNs) and LSTMs using PyTorch for time-series analysis of system performance data, achieving a 30% improvement in predictive maintenance accuracy.
- Deployed Big Data solutions using Apache Spark to process and analyze real-time streaming data, reducing latency in anomaly detection systems by 18%.
- Developed Generative Adversarial Networks (GANs) for synthetic data generation to augment limited datasets, enhancing model robustness for computer vision applications.
- Orchestrated containerized ML workflows using Docker and Kubernetes, enabling scalable model deployment across hybrid cloud environments on Azure.
- Applied Waterfall methodology to manage a structured development lifecycle for a customer-facing AI chatbot, ensuring on-time delivery within a 6-month timeline.
- Leveraged Apache Kafka to build real-time data pipelines for ingesting high-velocity sensor data, supporting predictive analytics for IoT based systems.
- Utilized Google Cloud Platform to deploy Transformers-based NLP models for sentiment analysis, improving customer feedback processing efficiency by 22%.
- Automated CI/CD pipelines using Jenkins to streamline model retraining and deployment, reducing deployment time by 35% for production-grade AI solutions.

EDUCATION

Master of Science in Artificial Intelligence Mar 2025 DePaul University, Chicago, United States

Bachelor of Technology in Electronics and Communications Aug 2021 HITAM (Hyderabad Institute of Technology and Management)